

COSC 3750

Scheduling

Kim Buckner

University of Wyoming

Mar. 01, 2022

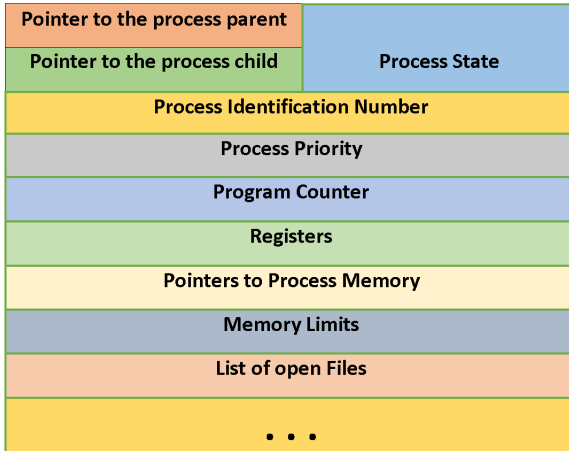
Processes

- A process is an active entity whereas a program is a passive one.
- Every process has a state (distinguished from 'state'):
 - new,
 - running,
 - waiting,
 - ready,
 - or terminated.

(more ...)

- Every process has a process control block (PCB) containing
 - its status,
 - pid,
 - PC,
 - registers,
 - memory limits,
 - open files,
 - memory map,
 - scheduling info,
 - accounting info, etc.

PCB



Created by NotesJam

CPU Scheduling

- Main idea in a multiprogramming environment: have some process running all the time.
- Process runs until it waits.
- Wait is usually for I/O.

Burst cycle

- Scheduling success depends a cycle—CPU time followed by I/O time
- Time is in “bursts.”
- This describes the duration of the portion of the cycle
- I/O bound process have very short (usually) CPU bursts and relatively long I/O bursts.
- Has been a lot of study, durations are usually exponential or hyper-exponential.

CPU Scheduler

- Short-term scheduler selects next process to run after CPU becomes idle.
- Ready queue is not necessarily FIFO.
- Queue generally contains PCBs.

Preemptive scheduling

- The decision to change a process state is made. Then ...
 - ① the process switches from running to waiting, or
 - ② the process switches from running to ready, or
 - ③ the process switches from waiting to ready, or
 - ④ the process terminates.

(more ...)

- states 1 & 4 require an immediate choice from the ready queue, easy decision.
- 1 & 4 = non-preemptive = Windows 3.1 and Apple Macintosh.
- Otherwise preemptive.
- Preemptive can be costly, i.e. shared data updates being interrupted.

(more ...)

- What happens when O/S is performing on behalf of process that is preempted?
 - Wait until it completes or
 - back off and cause it to be restarted when process runs again.
- Can of course just disable interrupts but this can be costly and may not scale well to multiprocessor systems

Dispatcher

- Gives control of CPU to a process.
- Switches context.
- Switches to user mode.
- Jumps to proper location to restart.
- Must be fast - dispatch latency can be a problem.

Scheduling criteria

- Lots of different methods have proposed and used.
- CPU utilization should range from 40 to 90 per cent but in real life the computation of this is often misleading.

Some bases (bā-sēz)

- Throughput - work being done
 - number of processes completing per time unit
- Turnaround time - time to run a process from submit to complete
 - sum of wait time, run time, and I/O time.
- Waiting time - process is runnable but not running.

(more ...)

- Response time - for interactive processes
 - from submit until responding starts but NOT time to output response
- Want to maximize CPU utilization and throughput and minimize waiting, turnaround and response time
- Some suggest minimizing variance in response time rather than reducing average

Scheduling algorithms

FCFS

- FCFS - First-come, First-served. Easily managed with a FIFO queue.
- Average wait time can be quite long.
- Generally not minimal, can be quite long, depends entirely on order of submission.
- Non-preemptive, bad for time-sharing systems.

SJF, shortest-job-first

- Based on last CPU burst.
- Provably optimal.
- But how to predict next burst?
- Usually with an exponential average.
 - $T_n + 1 = at_n + (1 - a)T_n$
 - Where T_n is the history and $0 \leq a \leq 1$ is the weight of the history and the last burst
 - often $a = 1/2$.
 - T_0 can be a constant or the system average, whatever

(SJF more . . .)

- Can be either preemptive or non-preemptive.
- Sometimes called
shortest-remaining-time-first if preemptive.

Priority scheduling

- Each process is associated with a priority.
- SJF is a form of this.
- Priorities are usually in a small fixed range:
 - 0-7,
 - 0-4095,
 - but no one can agree if 0 is the highest or lowest.

(more ...)

- We will use 'zero is low' unless we specify otherwise.
- Two types of criteria:
 - internal-time, cpu usage, number of open files,
 - external-importance to someone, like amount being paid for cpu use.
- Can be preemptive or non-preemptive.

(more ...)

- Main problem is starvation.
- Solution to that is aging, as a process stays in the run queue it keeps getting its priority improved until it finally runs

Round Robin

- Time-sharing systems.
- Preemption added to FCFS.
- Each process is allowed to run for a maximum of 1 time quantum.
- Ready queue is FIFO.
- Either the process will
 - run a quantum and be preempted or
 - it will make a request that causes it to relinquish the queue.

(more ...)

- Average wait time can be long.
- Each process gets $1/n$ of the CPU time and waits at most $(n - 1) * q$ time.
- If the quantum is large, RR becomes FCFS.
- One consideration is the time for a context switch.
- Want time quantum to be large with respect to the context switch.

(more ...)

- Turnaround time is affected as well but does not necessarily improve with an increase in quantum.
- Best if average CPU burst time is 80% of the time quantum.

Multi-level queue

- Different 'types' of processes have different priority queues.
- Each queue can have a different type of scheduling.
- Must be a scheduling between queues too, say preemptive priority, maybe with aging.

(more ...)

- Might could time slice the queues with each one having a separate time quantum.
- Or maybe just a percentage of CPU time for any particular queue.

M-L feedback

- Allows process to migrate between queues.
- Uses promotion and demotion.
- Each queue has a different scheduling method.
- etc.